

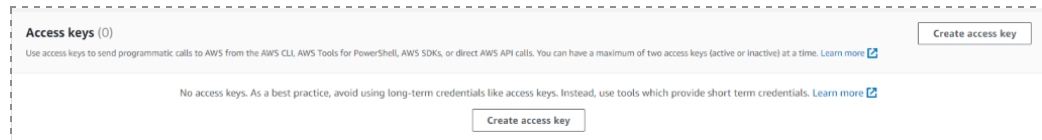
# Kubernetes and Kubeflow - Setup

---

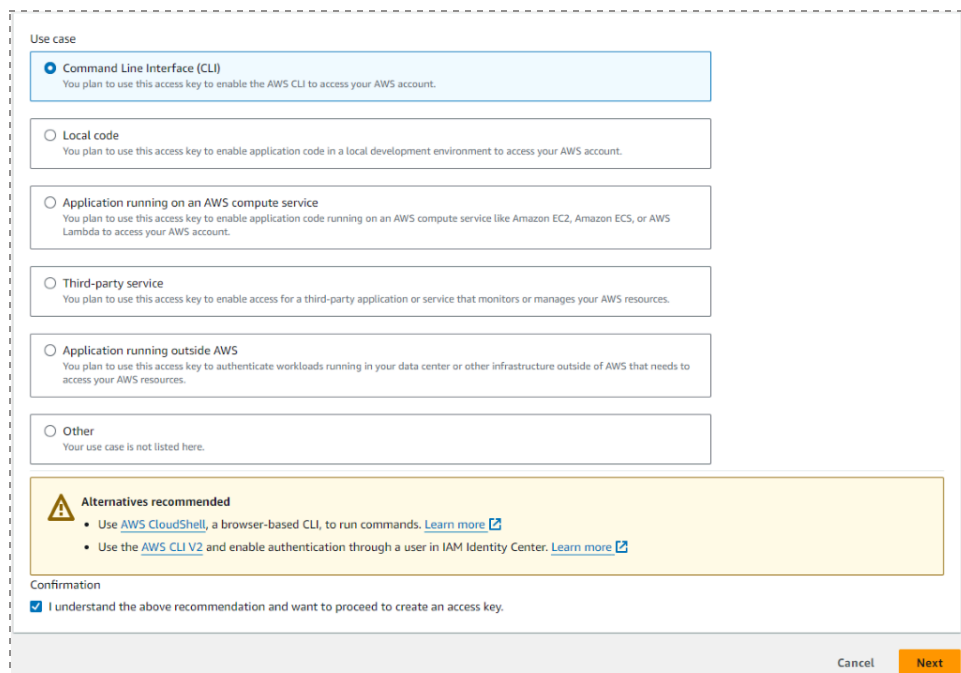
Prerequisites:

1. Get IAM access keys from AWS account

- You can access this by clicking on the top right your AWS account name and selecting **Security Credentials**
- Scroll to access keys and select **Create Access Key**



- A new screen will appear where the **Command Line Interface(CLI)** will be selected and check the confirmation box at the bottom and select next



- Add your description of the key and select **Create Access Key**
- It will then show your Access key and Secret access key. **STORE/SAVE THESE FOR LATER USE.** (Recommend saving to home directory)

2. Connect to instance shell through ssh

- Deactivate conda if it has not been already using: `deactivate conda`
- Note: if conda is activate (base) will show to the left of the user, you do not want that

# Preparing the EC2 instance (Create an EC2 instance if not already created)

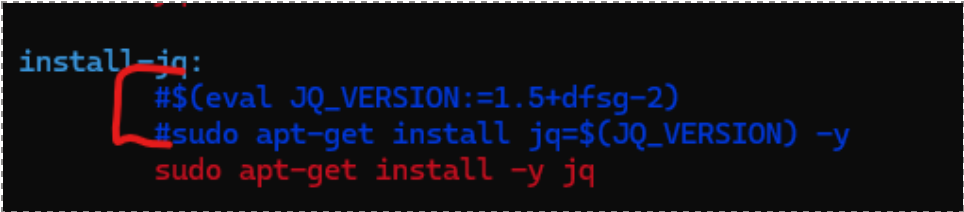
1. Run this command to install packages: `sudo apt update && sudo install -y git curl unzip tar make sudo vim wget`

2. Run these commands to install Kubeflow:

- `export KUBEFLOW_RELEASE_VERSION=v1.6.1`
- `export AWS_RELEASE_VERSION=v1.6.1-aws-b1.0.2`
- `git clone https://github.com/awslabs/kubeflow-manifests.git && cd kubeflow-manifests`
- `git checkout ${AWS_RELEASE_VERSION}`
- `git clone --branch ${KUBEFLOW_RELEASE_VERSION} https://github.com/kubeflow/manifests.git upstream`

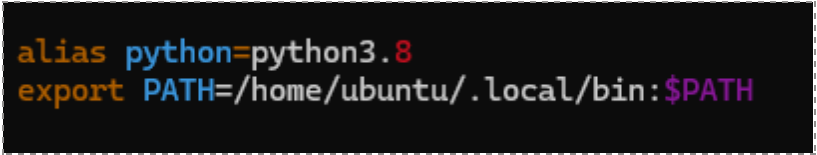
3. Edit the Configuration

- Run the command: `vim Makefile` (or `nano Makefile`, text editor doesn't matter, we just need to edit the Makefile)
- Scroll down to `install-jq` and comment out the two lines listed and Add the line: `sudo apt-get install -y jq` as shown below
- `***type "i" to insert text in vim***`



```
install-jq:
#$(eval JQ_VERSION:=1.5+dfsg-2)
#sudo apt-get install jq=${JQ_VERSION} -y
sudo apt-get install -y jq
```

- Save + Exit the Makefile (esc, then `":wq"`)
- Open the `bashrc` file using: `vim ~/.bashrc`
- Set the default version of python: `alias python=python3.8`
- Set the path variable..
- It should look like this:



```
alias python=python3.8
export PATH=/home/ubuntu/.local/bin:$PATH
```

- Exit + save the `~/.bashrc` file
- Run the command to reload the changes: `source ~/.bashrc`
- Deactivate conda if it is not already (command below): `conda deactivate`
- Run the automatic tool install process: `make install-tools`
- **\*If this command doesn't finish or gives dependency errors, anaconda needs to be removed using:** `rm -r /path/to/anaconda`

4. Configure AWS CLI

- **Run this to configure:** `aws configure --profile=kubeflow`
- AWS Access Key ID [None]: <enter your access key id>  
AWS Secret Access Key [None]: <enter your secret access key>
- Default region name [None]: us-east-1
- Default output format [None]: json
- Fill out the blanks (or stuff with <>) with the access key ID and secret access key obtained during prerequisite
- Edit the ~/.bashrc file again to include this on empty line: `export AWS_PROFILE=kubeflow`

```
export AWS_PROFILE=kubeflow
```

- Save + Exit the file
- Run the command to reload the changes: `source ~/.bashrc`
- Run this command to check if AWS CLI was configured correctly: `aws sts get-caller-identity`
- It should display an output that looks like this:
 

```
{
    "UserId": "<redacted>",

    "Account": "<redacted>",

    "Arn": "arn:aws:iam::<redacted>:user/<redacted>"
}
```

## 5. Create the cluster

- Run this to export the cluster name (Note: this name can be changed, this works for our use case):  
`export CLUSTER_NAME=kubeflow-2`  
`export CLUSTER_REGION=us-east-1`
- Run this to create and cluster and wait for it to be successfully created (20-30 minutes)  
`eksctl create cluster \`  
     `--name ${CLUSTER_NAME} \`  
     `--version 1.23 \`  
     `--region ${CLUSTER_REGION} \`  
     `--nodegroup-name demo-nodes \`  
     `--node-type m5.xlarge \`  
     `--nodes 5 \`  
     `--nodes-min 1 \`  
     `--nodes-max 10 \`  
     `--managed \`

--with-oidc

## 6. Install the Amazon EBS CSI Driver

- Create role by running: `eksctl create iamserviceaccount --name ebs-csi-controller-sa --namespace kube-system --cluster my-cluster --role-name AmazonEKS_EBS_CSI_DriverRole2 --role-only --attach-policy-arn arn:aws:iam::aws:policy/service-role/AmazonEBSCSIDriverPolicy --approve`
- Rule should show up in AWS console (not necessary in this case, we already create a rule that is added to the instances)
- For example, it would look like this:

<input type="checkbox"/>	<a href="#">AmazonEKS_EBS_CSI_DriverRole</a>	Identity Provider: arn:aws:iam::5332	-
<input type="checkbox"/>	<a href="#">AWSServiceRoleForAmazonEKS</a>	AWS Service: eks (Service-Linked Role)	Yesterday
<input type="checkbox"/>	<a href="#">AWSServiceRoleForAmazonEKSNodegroup</a>	AWS Service: eks-nodegroup (Service-Linked Role)	Yesterday
<input type="checkbox"/>	<a href="#">AWSServiceRoleForAutoScaling</a>	AWS Service: autoscaling (Service-Linked Role)	Yesterday
<input type="checkbox"/>	<a href="#">AWSServiceRoleForSupport</a>	AWS Service: support (Service-Linked Role)	-
<input type="checkbox"/>	<a href="#">AWSServiceRoleForTrustedAdvisor</a>	AWS Service: trustedadvisor (Service-Linked Role)	-
<input type="checkbox"/>	<a href="#">eksctl-kubeflow-1-addon-iamserviceaccount-kub-Role1-yWkIbr2zLdJ2</a>	Identity Provider: arn:aws:iam::5332	-
<input type="checkbox"/>	<a href="#">eksctl-kubeflow-1-cluster-ServiceRole-q8TKc0Ldvo6K</a>	AWS Service: eks	-
<input type="checkbox"/>	<a href="#">eksctl-kubeflow-1-nodegroup-demo-n-NodeInstanceRole-31eqCm8nwf6u</a>	AWS Service: ec2	-
<input type="checkbox"/>	<a href="#">kf-ack-sm-controller-role-kubeflow-1</a>	Identity Provider: arn:aws:iam::5332	-

- Create addon by replacing <accountid> with the account id of aws account as shown in photo (remove dashes in between):  
`eksctl create addon --name aws-ebs-csi-driver --cluster kubeflow-2 --service-account-role-arn arn:aws:iam::<accountid>:role/AmazonEKS_EBS_CSI_DriverRole2 -force`
- (change kubeflow-2 to your cluster name if changed)
- (change AmazonEKS\_EBS\_CSI\_DriverRole to the name of new IAM role if added)
- Get the update needed (change kubeflow-2 to your cluster name if changed):  
`eksctl get addon --name aws-ebs-csi-driver --cluster kubeflow-2`

(should list details, if nothing under update driver skip next step)

- Update addon to the version listed in the command output

*(if necessary change kubeflow-2 to cluster name)*

*(change <v1.31.0-eksbuild.1> to the correct updated driver without the <>)*

*(replace <accountid> with the account id as done above)*

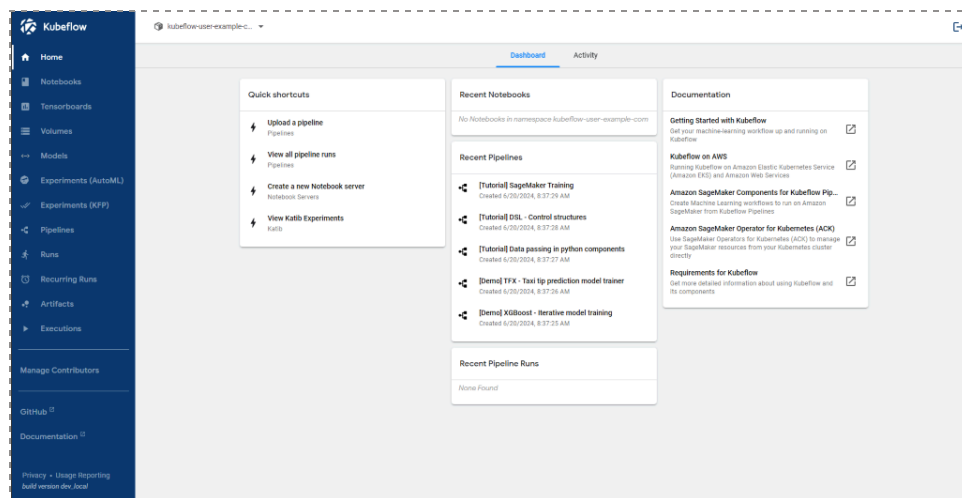
*(if necessary change AmazonEKS\_EBS\_CSI\_DriverRole to correct role created)*

```
eksctl update addon --name aws-ebs-csi-driver --version  
<v1.31.0-eksbuild.1> --cluster kubeflow-2 \
```

```
--service-account-role-arn
arn:aws:iam::<accountid>:role/AmazonEKS_EBS_CSI_DriverRole
2 -force
```

## 7. Add Kubeflow to the Cluster

- **Deploy Kubeflow:** `make deploy-kubeflow`  
`INSTALLATION_OPTION=kustomize DEPLOYMENT_OPTION=vanilla`
- **Verify all of the services are running in the pods:**  
`kubectl get pods -n cert-manager`  
`kubectl get pods -n cert-manager`  
`kubectl get pods -n auth`  
`kubectl get pods -n knative-eventing`  
`kubectl get pods -n knative-serving`  
`kubectl get pods -n kubeflow`  
`kubectl get pods -n kubeflow-user-example-com`
- If everything returns running, Kubeflow has been installed. If there any issues with specific services, try installing those services again
- Open up other terminal to allow dashboard connection (<desired\_local\_port> is the port you want to run locally):  
`ssh -i /path/to/identity_key.pem -L`  
`<desired_local_port>:127.0.0.1:8080`  
`ubuntu@<EC2_PUBLIC_IPV4_ADDRESS>`
- Check your directory with running: `pwd`
- If you are not already in `kubeflowmanifest`, run `cd kubeflowmanifest`
- Once you are in the correct directory, run: `make port-forward`
- Access Kubeflow from a browser using `127.0.0.1:<desired_local_port>`  
 Kubeflow's default username is: `user@example.com`  
 Kubeflow's default password is: `12341234`



## Scaling Nodes in Cluster:

Open a new terminal and run: `eksctl scale nodegroup --cluster=clusterName  
--nodes=desiredCount --name=nodegroupName`

*The Cluster must be scaled down to zero nodes after use, and then hibernated. In order for Kubeflow to run properly, it needs to be scaled up to five nodes.*